

CSV Data Quality CLI (Bash/Linux) — Terminal-Bench

One-pager for LinkedIn Media

What it does:

Small Bash CLI that analyzes two CSVs, normalizes CRLF→LF, applies deterministic sorting (LC_ALL=C), detects duplicates, and produces three reports.

How to run:

1) `chmod +x solution.sh`

2) `./solution.sh` # outputs `analysis_report.txt`, `data_summary.txt`, `duplicates_found.txt`

Outputs:

- `analysis_report.txt` — counts, basic file info, and tools used (`wc`, `sort`, `uniq`, `head`, `tail`).
- `data_summary.txt` — per-dataset totals and overall summary.
- `duplicates_found.txt` — per-file duplicate statistics (0 if none).

Engineering notes:

- Deterministic: `LC_ALL=C`; CRLF→LF normalization; set `-euo pipefail`.
- Offline test runner; no network installs in `run-tests.sh`.
- Docker-ready, pinned behavior; clear `--help/exit` codes (where applicable).

Results:

Local harness checks match assessment style: `oracle` = 100%, `nop` = 0%.

Stack:

Bash, Linux coreutils (`wc`, `sort`, `uniq`, `head`, `tail`), Docker, Git.

Links:

- Repository: <your GitHub link here>
- README includes examples and usage.